



Presentation Answers: Farm Financial RAG Application Analysis

1. How did you choose what to convert from the original DB schema to the new DB schema and why?

Schema Conversion Strategy: Original Source: The project started with a complex Excel file (`FINBIN Data Dictionary Farm.xlsx`) containing the original FINBIN (Farm Financial Benchmarking) database schema. Conversion Approach:

- **Simplified Complex Tables:** The original schema had extremely complex tables (e.g., `fm_guide` with 50+ columns). I simplified these to focus on core financial metrics that users would actually query.
- **Preserved Key Relationships:** Maintained the essential foreign key relationships between `hdb_main_data` (main farm records) and other tables.
- **Selected Essential Columns:** Instead of all 50+ columns in `fm_guide`, I kept only the most important financial ratios:
 - `current_ratio_beg/end` (liquidity)
 - `net_farm_income_cost/mkt` (profitability)
 - `working_capital_beg/end` (liquidity)
 - `debt_to_asset_ratios` (leverage)
 - `ebitda_cost/mkt` (operating performance)

Why This Approach:

- **LLM Token Efficiency:** Sending the entire 50+ column schema to the LLM would consume excessive tokens
- **Query Relevance:** Most user questions focus on basic financial metrics, not obscure technical fields
- **Maintainability:** A Simpler schema is easier to debug and extend
- **Performance:** Fewer columns = faster queries and better user experience

2. What are the pros and cons of this type of SQL building approach by the LLM?

PROS:

- **Natural Language Interface:** Users can ask questions in plain English without SQL knowledge

- Intelligent Query Generation: LLM understands context and generates appropriate JOINS, WHERE clauses, and aggregations
- Flexible Question Types: Handles complex analytical questions like "Which farms are in the top 10% for profitability?"
- Error Recovery: When queries fail, the LLM can suggest alternative phrasings
- Domain-Aware: The LLM understands farm financial terminology and relationships

CONS:

- Token Cost: Each query requires sending the entire schema to the LLM (expensive with large schemas)
- Inconsistent Results: LLM may generate different SQL for similar questions
- Security Risks: SQL injection potential if not properly sanitized
- Performance Overhead: LLM call adds 1-3 seconds to each query
- Limited Complex Logic: Struggles with very complex analytical queries requiring multiple CTEs or window functions
- Schema Dependency: Changes to the database structure require updating prompts

Mitigation Strategies Used:

- Simplified schema to reduce token usage
- Added explicit table/column lists in prompts
- Implemented query validation and error handling
- Used temperature=0.1 for more consistent results

3. What would happen if there were more than three rows of data?

Current State: The application is designed with 3 sample farms across 10 tables (30 total records). With More Data (1000+ farms): Performance Impact:

- Query Speed: SQLite can handle millions of rows efficiently - minimal impact on query execution
- LLM Processing: No effect on LLM response time (schema size matters more than data size)
- Memory Usage: Minimal increase as only result previews are loaded into memory

Scalability Considerations:

- Database Size: 1000 farms \approx 10,000 records = ~50MB database (very manageable)
- Query Optimization: Would need indexes on frequently queried columns (state, year, financial metrics)
- Result Pagination: Currently limited to 10-20 rows in previews - would need pagination for large result sets
- Caching: Could implement query result caching for frequently asked questions

Enhanced Capabilities:

- Better Analytics: More data enables meaningful statistical analysis and benchmarking
- Trend Analysis: Sufficient data for year-over-year comparisons
- Geographic Insights: Better state/county performance comparisons
- Percentile Rankings: More accurate quartile and percentile calculations

Code Changes Needed:

4. How do you choose what questions to use for testing?

Testing Question Selection Strategy:

1. Progressive Complexity:

python

```
# From demo.py and test files

demo_questions = [

    "How many farms are in the database?",          # Basic count

    "Which farms have the highest current ratio?",  # Simple ranking

    "What is the average working capital by state?", # Aggregation + grouping

    "Show me farms with debt-to-equity ratio above 2.0", # Filtering

    "How did net worth change from the beginning to end of year?" # Complex analysis

]
```

2. Category-Based Testing:

- Financial Performance: Ratios, rankings, profitability metrics
- Geographic Analysis: State/county comparisons, regional trends
- Trends and Changes: Time-based analysis, year-over-year comparisons
- Benchmarking: Percentiles, quartiles, top/bottom performers

3. Edge Case Testing:

- Empty Results: "Show me farms with negative working capital"
- Error Handling: "What's the average of a non-existent column?"
- Complex Joins: Questions requiring multiple table relationships
- Aggregation Limits: "What's the 99th percentile for current ratio?"

4. User Experience Testing:

- Natural Language Variations: "Which farms are most profitable?" vs "Show me farms with the highest net income"
- Ambiguous Questions: "Best farms" (requires clarification)
- Follow-up Questions: Testing context retention

5. Performance Testing:

```
# From test_rag_app.py

test_questions = [

    "How many farms are in the database?",      # Fast count query

    "What is the average working capital?",     # Aggregation

    "Which state has the most farms?",          # Grouping

    "Show me farms with the highest current ratio" # Sorting

]
```

5. How do you track what is happening through the multiple steps?

Multi-Step Tracking Architecture. Structured Logging System:

python

```
# From farm_rag_app.py

import logging

logging.basicConfig(level=logging.INFO)

logger = logging.getLogger(__name__)

def ask_question(self, user_question: str) -> Dict[str, Any]:

    try:

        logger.info(f"Processing question: {user_question}")

        # Step 1: Generate SQL query

        sql_query = self._generate_sql_query(user_question)

        # Step 2: Execute SQL query

        query_result = self._execute_sql_query(sql_query)

        # Step 3: Generate natural language response

        response = self._generate_response(user_question, query_result)
```

2. Comprehensive Result Tracking:

python

```

@dataclass
class QueryResult:
    success: bool
    data: Optional[pd.DataFrame]
    sql_query: str
    error_message: Optional[str] = None
    row_count: int = 0
    execution_time: float = 0.0 # Performance tracking

```

3. Step-by-Step Metadata:

python

```

result = {
    "success": True,
    "question": user_question,
    "sql_query": sql_query, # Step 1 output
    "response": response, # Step 3 output
    "query_result": {
        "success": query_result.success, # Step 2 status
        "row_count": query_result.row_count, # Step 2 results
        "execution_time": query_result.execution_time, # Performance
        "error_message": query_result.error_message
    }
}

```

4. API Response Tracking:

python

```

# From farm_rag_api.py
class QuestionResponse(BaseModel):
    success: bool
    question: str
    sql_query: str
    response: str

```

```
query_result: Dict[str, Any]

data_preview: Optional[list] = None

error: Optional[str] = None
```

5. Testing and Validation Tracking:

python

```
# From test_rag_app.py - Comprehensive test suite

tests = [

    ("Environment Configuration", test_environment),

    ("OpenAI API Connection", test_openai_connection),

    ("Database Connection", test_database),

    ("RAG Application Import", test_rag_import),

    ("RAG Instance Creation", test_rag_instance),

    ("SQL Generation", test_sql_generation),      # Step 1 tracking

    ("SQL Execution", test_sql_execution),        # Step 2 tracking

    ("Complete RAG Workflow", test_full_rag_workflow), # End-to-end tracking

    ("API Endpoints", test_api_endpoints),

    ("Performance Test", run_performance_test)

]
```

6. Performance Monitoring:

python

```
# Timing each step

start_time = time.time()

result = rag_app.ask_question(question)

end_time = time.time()

response_time = end_time - start_time

# Track in results

results.append({

    "question": question,

    "success": True,

    "response_time": response_time,
```

```
        "rows_returned": result["query_result"]["row_count"]
    })
```

7. Error Tracking and Recovery:

python

```
try:
    # Process question
    result = rag_app.ask_question(question)

except Exception as e:
    logger.error(f"Error in ask_question: {e}")

    return {
        "success": False,
        "question": user_question,
        "error": str(e),
        "response": f"I apologize, but I encountered an error: {e}"
    }
```

Key Tracking Benefits:

- Debugging: Easy to identify which step failed
- Performance Analysis: Track bottlenecks in the pipeline
- User Experience: Provide detailed feedback on query execution
- Monitoring: Log all interactions for analysis and improvement
- Testing: Comprehensive validation of each component

This multi-step tracking system provides complete visibility into the RAG pipeline, enabling easy debugging of issues, performance optimization, and delivering detailed feedback to users regarding their queries.

ER Diagram Summary

Core Tables:

1. hdb_main_data (Main Entity)

- Primary Key: `hdb_main_data_id`
 - Purpose: Central farm record with identification, location, and metadata
 - Key Fields: Farm ID, state, county, client name, year, analyst info
1. `fm_genin` (Farm General Information)
 - Primary Key: `fm_genin_guid`
 - Foreign Key: `hdb_main_data_id` → `hdb_main_data`
 - Purpose: Farm name and general information linking
 1. `fm_guide` (Financial Guide - Core Financial Metrics)
 - Primary Key: `id` (auto-increment)
 - Foreign Keys: `fm_genin_guid`, `hdb_main_data_id`
 - Purpose: Key financial ratios and performance metrics
 - Key Metrics: Current ratio, working capital, net farm income, debt ratios, EBITDA
 1. `fm_stmts` (Financial Statements)
 - Primary Key: `id` (auto-increment)
 - Foreign Keys: `fm_genin_guid`, `hdb_main_data_id`
 - Purpose: Detailed financial statement data
 - Key Fields: Net worth, cash flow, income statements, balance sheet changes

Supporting Tables:

- `fm_prf_lq`: Profitability & Liquidity analysis
- `fm_cap_ad`: Capital & Asset data
- `fm_hhold`: Household financial data
- `fm_nf_ie`: Non-farm income & expenses
- `fm_fm_exp`: Farm expenses
- `fm_fm_inc`: Farm income
- `fm_beg_bs_end_bs`: Beginning/Ending balance sheet data

Relationship Pattern:

- One-to-Many: Each farm (`hdb_main_data`) can have multiple records in each financial table
- Dual Foreign Keys: Most tables reference both `hdb_main_data_id` and `fm_genin_guid` for data integrity
- Hierarchical Structure: `hdb_main_data` → `fm_genin` → Financial tables

Key Design Features:

- Simplified Schema: Focused on essential financial metrics for RAG queries
- Flexible Structure: Supporting tables can be extended with additional columns
- Data Integrity: Foreign key constraints ensure referential integrity
- Query Optimization: Primary keys and foreign keys enable efficient JOINS

This ER diagram represents the simplified, RAG-optimized version of the original FINBIN database schema, designed for efficient natural language querying and analysis.

[Review Changes](#)

Detailed Component Interactions

1. User Interface Layer

- Web Interface (`web_interface.html`): Modern HTML5 interface with JavaScript
- CLI Interface (`farm_rag_app.py`): Direct command-line interaction
- Demo Interface (`demo.py`): Guided demonstration with example questions
- Test Interface (`quick_test.py`): Quick validation and testing

2. API Layer (FastAPI)

- Main API Server (`farm_rag_api.py`): RESTful API endpoints
- CORS Middleware: Cross-origin resource sharing
- Static File Server: Serves HTML, CSS, JS files
- Health Monitoring: System status and diagnostics

3. Core RAG Engine

- FarmDataRAG Class: Main orchestrator
- SQL Generator: Converts natural language to SQL using OpenAI
- SQL Executor: Executes queries against SQLite database
- Response Generator: Creates natural language responses using OpenAI

4. External Services

- OpenAI API: GPT models for SQL generation and response creation
- Environment Variables: Configuration management (API keys, settings)

5. Database Layer

- SQLite Database: Local file-based database
- Schema: 11 tables with a financial data structure
- Sample Data: 10 farms with comprehensive financial records

6. Database Management

- Database Creator: Initial schema and data setup
- Data Adder: Additional sample data insertion
- Simple DB Creator: Windows-optimized database creation
- Database Checker: Validation and health monitoring

7. Testing & Validation

- Comprehensive Tests: Full system testing suite
- Performance Tests: Load and response time testing
- Health Checks: System status monitoring

8. Deployment Scripts

- Startup Script: Linux/macOS automated startup
- Windows Batch: Windows user setup automation
- Windows PowerShell: Advanced Windows automation



Data Flow Process

Question Processing Flow:

1. User Input → Web Interface/CLI
1. HTTP Request → FastAPI Server
1. Question Processing → FarmDataRAG Class
1. SQL Generation → OpenAI API (GPT-3.5/GPT-4)
1. Query Execution → SQLite Database
1. Response Generation → OpenAI API
1. Result Formatting → FastAPI Response
1. Display → Web Interface/CLI

Database Management Flow:

1. Schema Creation → Database Creator
1. Data Population → Data Adder
1. Validation → Database Checker
1. Query Processing → RAG Engine

Testing Flow:

1. Environment Check → Quick Test
1. Component Testing → Comprehensive Tests
1. Performance Testing → Performance Tests
1. Health Monitoring → Health Checks

Deployment Options

Development Mode:

- Direct Python execution
- Individual component testing
- Interactive debugging

Production Mode:

- FastAPI server with Uvicorn
- Web interface access
- API documentation
- Health monitoring

Windows Deployment:

- Automated batch scripts
- PowerShell automation
- User-friendly setup process

This architecture provides a robust, scalable, and maintainable RAG application with multiple interfaces, comprehensive testing, and cross-platform deployment capabilities.